

Week 1 Exercise

INTERNET PRIVACY POLL

Internet privacy has gained widespread attention in recent years. To measure the degree to which people are concerned about hot-button issues like Internet privacy, social scientists conduct polls in which they interview a large number of people about the topic. In this assignment, we will analyze data from a July 2013 Pew Internet and American Life Project poll on Internet anonymity and privacy, which involved interviews across the United States. While the full polling data can be found [here](#), we will use a more limited version of the results, available in [AnonymityPoll.csv](#). The dataset has the following fields (all Internet use-related fields were only collected from interviewees who either use the Internet or have a smartphone):

- **Internet.Use:** A binary variable indicating if the interviewee uses the Internet, at least occasionally (equals 1 if the interviewee uses the Internet, and equals 0 if the interviewee does not use the Internet).
- **Smartphone:** A binary variable indicating if the interviewee has a smartphone (equals 1 if they do have a smartphone, and equals 0 if they don't have a smartphone).
- **Sex:** Male or Female.
- **Age:** Age in years.
- **State:** State of residence of the interviewee.
- **Region:** Census region of the interviewee (Midwest, Northeast, South, or West).
- **Conservativeness:** Self-described level of conservativeness of interviewee, from 1 (very liberal) to 5 (very conservative).
- **Info.On.Internet:** Number of the following items this interviewee believes to be available on the Internet for others to see: (1) Their email address; (2) Their home address; (3) Their home phone number; (4) Their cell phone number; (5) The employer/company they work for; (6) Their political party or political affiliation; (7) Things they've written that have their name on it; (8) A photo of them; (9) A video of them; (10) Which groups or organizations they belong to; and (11) Their birth date.
- **Worry>About.Info:** A binary variable indicating if the interviewee worries about how much information is available about them on the Internet (equals 1 if they worry, and equals 0 if they don't worry).
- **Privacy.Importance:** A score from 0 (privacy is not too important) to 100 (privacy is very important), which combines the degree to which they find privacy important in the following: (1) The websites they browse; (2) Knowledge of the place they

are located when they use the Internet; (3) The content and files they download; (4) The times of day they are online; (5) The applications or programs they use; (6) The searches they perform; (7) The content of their email; (8) The people they exchange email with; and (9) The content of their online chats or hangouts with others.

- **Anonymity.Possible:** A binary variable indicating if the interviewee thinks it's possible to use the Internet anonymously, meaning in such a way that online activities can't be traced back to them (equals 1 if he/she believes you can, and equals 0 if he/she believes you can't).

- **Tried.Masking.Identity:** A binary variable indicating if the interviewee has ever tried to mask his/her identity when using the Internet (equals 1 if he/she has tried to mask his/her identity, and equals 0 if he/she has not tried to mask his/her identity).

- **Privacy.Laws.Effective:** A binary variable indicating if the interviewee believes United States law provides reasonable privacy protection for Internet users (equals 1 if he/she believes it does, and equals 0 if he/she believes it doesn't).

PROBLEM 1.1 - LOADING AND SUMMARIZING THE DATASET

Using `read.csv()`, load the dataset from [AnonymityPoll.csv](#) into a data frame called `poll` and summarize it with the `summary()` and `str()` functions.

How many people participated in the poll?

Answer:

PROBLEM 1.2 - LOADING AND SUMMARIZING THE DATASET

Let's look at the breakdown of the number of people with smartphones using the `table()` and `summary()` commands on the `Smartphone` variable. (HINT: These three numbers should sum to 1002.)

How many interviewees responded that they use a smartphone?

Answer:

How many interviewees responded that they don't use a smartphone?

Answer:

How many interviewees did not respond to the question, resulting in a missing value, or NA, in the `summary()` output?

Answer:

PROBLEM 1.3 - LOADING AND SUMMARIZING THE DATASET

By using the `table()` function on two variables, we can tell how they are related. To use the `table()` function on two variables, just put the two variable names inside the parentheses, separated by a comma (don't forget to add `poll$` before each variable name). In the output, the possible values of the first variable will be listed in the left, and the possible values of the second variable will be listed on the top. Each entry of the table counts the number of observations in the data set that have the value of the first value in that row, and the value of the second variable in that column. For example, suppose we want to create a table of the variables

"Sex" and "Region". We would type

```
table(poll$Sex, poll$Region)
```

in our R Console, and we would get as output

```
Midwest Northeast South West
```

```
Female 123 90 176 116
```

```
Male 116 76 183 122
```

This table tells us that we have 123 people in our dataset who are female and from the Midwest, 116 people in our dataset who are male and from the Midwest, 90 people in our dataset who are female and from the Northeast, etc.

You might find it helpful to use the `table()` function to answer the following questions:

Which of the following are states in the Midwest census region? (Select all that apply.)

Colorado

Kansas

Kentucky

Missouri

Ohio

Pennsylvania

Which was the state in the South census region with the largest number of interviewees?

Alabama

Arkansas

Delaware
Florida
Georgia
Kentucky
Louisiana
Maryland
Mississippi
South Carolina
Oklahoma
South Carolina
Tennessee
Texas
Virginia
West Virginia

PROBLEM 2.1 - INTERNET AND SMARTPHONE USERS

As mentioned in the introduction to this problem, many of the response variables (Info.On.Internet, [Worry.About.Info](#), Privacy.Importance, Anonymity.Possible, and Tried.Masking.Identity) were not collected if an interviewee does not use the Internet or a smartphone, meaning the variables will have missing values for these interviewees.

How many interviewees reported not having used the Internet and not having used a smartphone?

Answer:

How many interviewees reported having used the Internet and having used a smartphone?

Answer:

How many interviewees reported having used the Internet but not having used a smartphone?

Answer:

How many interviewees reported having used a smartphone but not having used the Internet?

Answer:

PROBLEM 2.2 - INTERNET AND SMARTPHONE USERS

How many interviewees have a missing value for their Internet use?

Answer:

How many interviewees have a missing value for their smartphone use?

Answer:

PROBLEM 2.3 - INTERNET AND SMARTPHONE USERS

Use the subset function to obtain a data frame called "limited", which is limited to interviewees who reported Internet use or who reported smartphone use. In lecture, we used the & symbol to use two criteria to make a subset of the data. To only take observations that have a certain value in one variable or the other, the | character can be used in place of the & symbol. This is also called a logical "or" operation.

How many interviewees are in the new data frame?

Answer:

Important: For all remaining questions in this assignment please use the limited data frame you created in Problem 2.3.

PROBLEM 3.1 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

Which variables have missing values in the limited data frame? (Select all that apply.)

Internet.Use
Smartphone
Sex
Age
State
Region
Conservativeness
Info.On.Internet
[Worry.About.Info](#)

Privacy.Importance
Anonymity.Possible
Tried.Masking.Identity
Privacy.Laws.Effective

PROBLEM 3.2 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

What is the average number of pieces of personal information on the Internet, according to the Info.On.Internet variable?

Answer:

PROBLEM 3.3 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

How many interviewees reported a value of 0 for Info.On.Internet?

Answer:

How many interviewees reported the maximum value of 11 for Info.On.Internet?

Answer:

PROBLEM 3.4 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

What proportion of interviewees who answered the [Worry>About.Info](#) question worry about how much information is available about them on the Internet? Note that to compute this proportion you will be dividing by the number of people who answered the [Worry>About.Info](#) question, not the total number of people in the data frame.

Answer:

PROBLEM 3.5 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

What proportion of interviewees who answered the Anonymity.Possible question think it is possible to be completely anonymous on the Internet?

Answer:

PROBLEM 3.6 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

What proportion of interviewees who answered the Tried.Masking.Identity question have tried masking their identity on the Internet?

Answer:

PROBLEM 3.7 - SUMMARIZING OPINIONS ABOUT INTERNET PRIVACY

What proportion of interviewees who answered the Privacy.Laws.Effective question find United States privacy laws effective?

Answer:

PROBLEM 4.1 - RELATING DEMOGRAPHICS TO POLLING RESULTS

Often, we are interested in whether certain characteristics of interviewees (e.g. their age or political opinions) affect their opinions on the topic of the poll (in this case, opinions on privacy). In this section, we will investigate the relationship between the characteristics Age and Smartphone and outcome variables Info.On.Internet and Tried.Masking.Identity, again using the limited data frame we built in an earlier section of this problem.

Build a histogram of the age of interviewees.

What is the best represented age group in the population?

People aged about 20 years old

People aged about 40 years old

People aged about 60 years old

People aged about 80 years old

PROBLEM 4.2 - RELATING DEMOGRAPHICS TO POLLING

RESULTS

Both Age and Info.On.Internet are variables that take on many values, so a good way to observe their relationship is through a graph. We learned in lecture that we can plot Age against Info.On.Internet with the command `plot(limited$Age, limited$Info.On.Internet)`. However, because Info.On.Internet takes on a small number of values, multiple points can be plotted in exactly the same location on this graph.

What is the largest number of interviewees that have exactly the same value in their Age variable AND the same value in their Info.On.Internet variable? In other words, what is the largest number of overlapping points in the plot `plot(limited$Age, limited$Info.On.Internet)`? (HINT: Use the table function to compare the number of observations with different values of Age and Info.On.Internet.)

Answer:

PROBLEM 4.3 - RELATING DEMOGRAPHICS TO POLLING RESULTS

To avoid points covering each other up, we can use the `jitter()` function on the values we pass to the plot function. Experimenting with the command `jitter(c(1, 2, 3))`, what appears to be the functionality of the jitter command?

`jitter` randomly reorders the values passed to it, and two runs will yield the same result

`jitter` randomly reorders the values passed to it, and two runs will yield different results

`jitter` adds or subtracts a small amount of random noise to the values passed to it, and two runs will yield the same result

`jitter` adds or subtracts a small amount of random noise to the values passed to it, and two runs will yield different results

Answer:

PROBLEM 4.4 - RELATING DEMOGRAPHICS TO POLLING RESULTS

Now, plot Age against Info.On.Internet with `plot(jitter(limited$Age), jitter(limited$Info.On.Internet))`. What relationship do you observe between Age and

Info.On.Internet?

Older age seems strongly associated with a larger value for Info.On.Internet
Older age seems moderately associated with a larger value for Info.On.Internet
Older age does not seem associated with a change in the value of Info.On.Internet
Older age seems moderately associated with a smaller value for Info.On.Internet
Older age seems strongly associated with a smaller value for Info.On.Internet

Answer:

PROBLEM 4.5 - RELATING DEMOGRAPHICS TO POLLING RESULTS

Use the `tapply()` function to obtain the summary of the `Info.On.Internet` value, broken down by whether an interviewee is a smartphone user.
What is the average `Info.On.Internet` value for smartphone users?

Answer:

What is the average `Info.On.Internet` value for non-smartphone users?

PROBLEM 4.6 - RELATING DEMOGRAPHICS TO POLLING RESULTS

Similarly use `tapply` to break down the `Tried.Masking.Identity` variable for smartphone and non-smartphone users.

What proportion of smartphone users who answered the `Tried.Masking.Identity` question have tried masking their identity when using the Internet?

Answer:

What proportion of non-smartphone users who answered the `Tried.Masking.Identity` question have tried masking their identity when using the Internet?

Answer:

